

Reconfigurable Systems



Jorge Semião

Instituto Superior de Engenharia

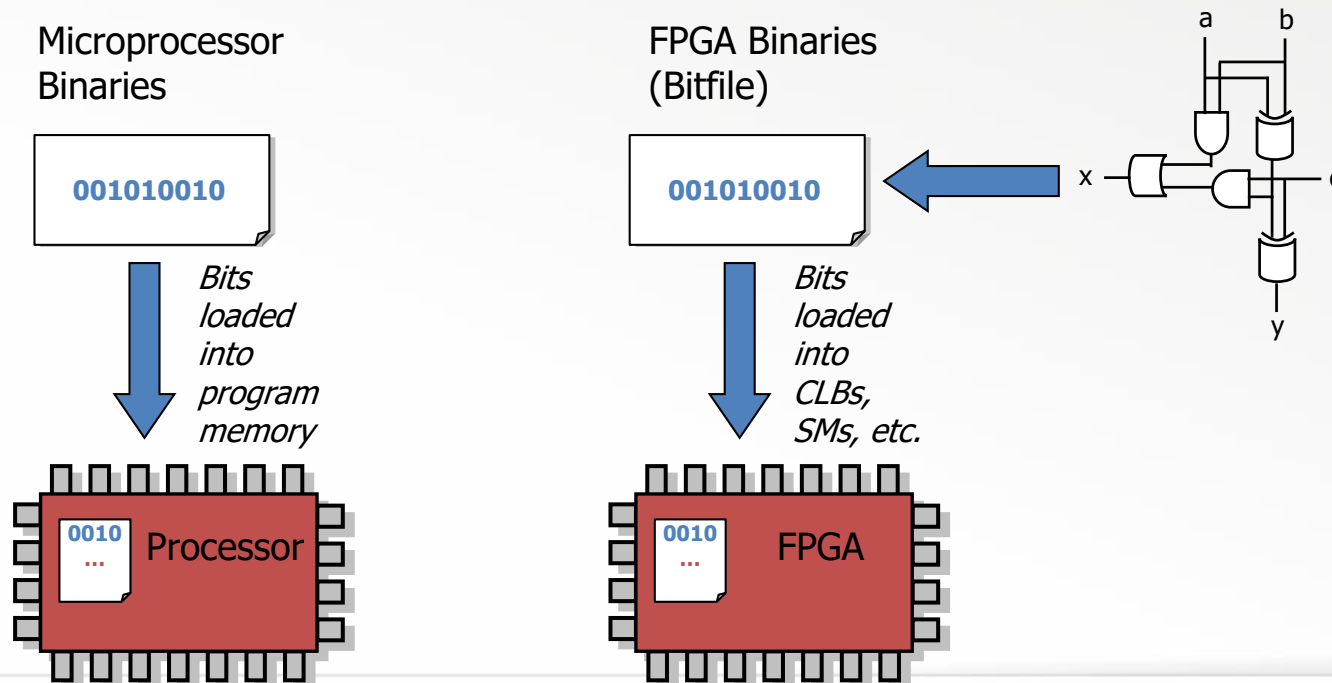
www.ualg.pt

Introduction to Reconfigurable Systems

- Can also be referred to Reconfigurable Computing
- Can be thought of as software-defined functionality, where flexibility is controlled predominately through the specification of bit patterns.
- Is the study of architectures that can adapt (after fabrication) to a specific application or application domain
 - Involves architecture, design strategies, tool flows, CAD, languages, algorithms, programmable hardware

What is Reconfigurable Computing?

- Alternatively, RS/RC is a way of implementing circuits without fabricating a device
 - Essentially allows circuits to be implemented as “software”
 - “circuits” are no longer the same thing as “hardware”
 - RC devices are programmable by downloading bits - just like software



Why is RS/RC important?

- Tremendous performance advantages
 - In some cases, > 100x faster than microprocessor
 - Alternatively, similar performances as large cluster
 - But smaller, lower power, cheaper, etc.
 - Example:

```
for (i=0; i < 16; i++)  
  y += c[i] * x[i]
```
- Software executes sequentially
- RC executes all multiplications in parallel
 - Additions become tree of adders
- Even with slower clock, RC is likely much faster
- Performance difference even greater for larger input sizes
 - SW time increases linearly - $O(n)$
 - RC time is basically $O(\log_2(n))$ - If enough area is available

When to use RS/RC?

Implementation Possibilities



Why not use an ASIC for everything?

Moore's Law

- Moore's Law is the empirical observation made in 1965 that the number of transistors on an integrated circuit doubles every 18 months [Wikipedia]

■ 1993: 1 Million transistors

2007: >1 BILLION
transistors!!!!

Becoming
extremely difficult
to design this -
ASICs are
expensive!

Moore's Law

- Solution: Make billions of transistors into a reconfigurable fabric - fabricate 1 big chip and use it for many things
 - Area overhead: circuit in FPGA can require 20x more transistors
 - But, that's still equivalent to a > 50 million transistor ASIC
 - Pentium IV ~ 42 million transistors
 - Modern FPGAs reportedly support millions of logic gates!

2007: >1 BILLION
transistors!!!!

Solution: Make this
reconfigurable

When should RS/RC be used?

- 1) When it provides the cheapest solution
 - Depends on:
 - NRE Cost - Non-recurring engineering cost
 - Cost involved with designing system
 - Unit cost - cost of a manufacturing/purchasing a single device
 - Volume - # of units
 - Total cost = $\text{NRE} + \text{unit cost} * \text{volume}$
 - RC is typically more cost effective for low volume devices
 - RC: low NRE, high unit cost
 - ASIC: very high NRE, low unit cost

What about microprocessors?


- Similar cost issues
 - uPs
 - low NRE cost (coding is cheap)
 - Unit cost varies from several dollars to several thousand
- Wouldn't cheapest microprocessor always be the cheapest solution?
 - Yes, but ...

What about microprocessors?

- Often, microprocessors cannot meet performance constraints
 - e.g. video decoder must achieve minimum frame rate
 - Common reason for using custom circuit implementation

Example

- FPGA: Unit cost = 5, NRE cost = 200,000
- Microprocessor (μ P): Unit cost = 8, NRE cost = 100,000
- Problem: Find cheapest implementation for all possible volumes (assume both implementations meet constraints)



$$5v + 200k = 8v + 100k$$
$$v = 33k$$

Answer: For volumes less than 33k, μ P is cheapest solution. For all other volumes, FPGA is cheapest solution.

Example: Your Turn


- FPGA
 - Unit cost: 6, NRE cost: 300,000
- ASIC
 - Unit cost: 2, NRE cost: 3,000,000
- Microprocessor (μ P)
 - Unit cost: 10, NRE cost: 100,000
- Problem: Find cheapest implementation for all possible volumes (assume that all possibilities meet performance constraints)

Note: NRE meaning

Non-recurring Engineering. Non-recurring engineering cost refers to the one-time cost to research, design, develop and test a new product or product enhancement. When budgeting for a new product, NRE must be considered to analyze if a new product will be profitable. [Wikipedia]

Another Example

- FPGA
 - Unit cost: 7, NRE cost: 300,000
- ASIC
 - Unit cost: 4, NRE cost: 3,000,000
- Microprocessor (μ P)
 - Unit cost: 1, NRE cost: 100,000




*Answer: μ P cheapest solution
at any volume – not
uncommon*

When should RS/RC be used?

- 2) When time to market is critical
 - Huge effect on total revenue

RC has faster time to market than ASIC



When should RS/RC be used?

- 3) When circuit may have to be modified
 - Can't change ASIC - hardware
 - Can change circuit implemented in FPGA
- Uses
 - When standards change
 - Codec changes after devices fabricated
 - Allows addition of new features to existing devices
 - Fault tolerance/recovery
 - “Partial reconfiguration” allows virtual fabric size - analogous to virtual memory
- Without RS/RC
 - Anything that may have to be reconfigured is implemented in software
 - Performance loss

Design Space Exploration

1. Determine architectures that meet performance requirements
 - Not trivial, requires performance analysis/estimation - important problem
 - Will study later in semester
 - And, other constraints - power, size, etc.
 2. Estimate volume of device
 3. Determine cheapest solution
- The best architecture for an application is typically the cheapest one that meets all design constraints.



RS/RC Markets

- Embedded Systems
 - FPGAs appearing in set-top boxes, routers, audio equipment, etc.
 - Advantages
 - RC achieves performance close to ASIC, sometimes at much lower cost
 - Many other embedded systems still use ASIC due to high volume
 - » Cell phones, iPod, game consoles, etc.
 - Reconfigurable!
 - If standards changes, architecture is not fixed
 - Can add new features after production



RS/RC Markets

- High-performance embedded computing (HPEC)
 - High-performance/super computing with special needs (low power, low size/weight, etc.)
 - Satellite image processing
 - Target recognition
 - RS/RC Advantages
 - Much smaller/lower power than a supercomputer
 - Fault tolerance



RS/RC Markets

- High-performance computing - HPC
 - Cray XD-1
 - 12 AMD Opterons, FPGAs
 - SGI Altix
 - 64 Itaniums, FPGAs
 - IBM Chameleon
 - Cell processor, FPGAs
 - Many others
- RS/RC advantages
 - HPC used for many scientific apps
 - Low volume, ASIC rarely feasible




RS/RC Markets

- General-purpose computing???
 - Ideal situation: desktop machine/OS uses RC to speedup up all applications
 - Problems
 - RS/RC can be very fast, but not for all applications
 - Generally requires parallel algorithms
 - Coding constructs used in many applications not appropriate for hardware
 - Subject of tremendous amount of past and likely future research
- How to use extra transistors on general purpose CPUs?
 - More cache
 - More microprocessors
 - FPGA
 - Something else?


Limitations of RS/RC

- 1) Not all applications can be improved

Embedded Applications – Large Speedups



Desktop Applications – *No Speedup*



- 2) Tools need serious improvement!
- 3) Design strategies are often ad-hoc
- 4) Floating point?
 - Requires a lot of area, but becoming practical